

# NMF Final (Only nndsvd 5 component without ozone)

William Zhang, Eva, Jerry, Meredith

2025-01-24

```
# load the packages
library(NMF)
library(tidyverse)
library(grid)
library(gridExtra)
library(readxl)
library(circular)
library(lwgeom)
library(units)
```

## Procedure

1. Remove hourly observation with missing observation for any chemical
2. Remove background noise level using min values (except for chemicals with minimum value  $< 2 \times \text{LOD}$  and maximum value  $> 100 \times \text{LOD}$ )
3. Zero values are converted to a random value between 0 and  $0.5 \times \text{LOD}$
4. Normalize using min and max
5. Remove Ozone (wouldn't affect # of obs.)

## Reading the data

```
hourly_data <- readRDS("../DataProcessing/Trailer_hourly_merge_20240905.rds")

# PROCEDURE STEP 1:
hourly_data <- hourly_data %>% rename('co2' = 'co2_ppm')

vocs <- c("ethane", "ethene", "propane", "propene",
          "1_3-butadiene", "i-butane", "n-butane",
          "acetylene", "cyclopentane", "i-pentane",
          "n-pentane", "n-hexane", "isoprene", "n-heptane",
          "benzene", "n-octane", "toluene", "ethyl-benzene",
          "m&p-xylene", "o-xylene")

non_vocs <- c('ch4', 'co2', 'co', 'h2s', 'so2', 'nox', 'o3')

# remove row with missing obs for any chemical
hourly_nona <- hourly_data %>%
  select(any_of(c('day', 'time_utc', vocs, non_vocs, 'wdr_deg', 'wsp_ms')))) %>%
  na.omit()
```

```

# retrieving the vocs, removing everything else except the vocs
hourly_vocs <- hourly_nona %>% select(any_of(vocs))

# retrieving the non-vocs: co2_ppm, nox, ch4, h2s, so2, o3
# double check this
hourly_non_vocs <- hourly_nona %>% select(any_of(non_vocs))


hourly_full_nona <- cbind(hourly_non_vocs, hourly_vocs)

# retrieve a vector of yearmonth
hourly_dates <- hourly_nona %>%
  mutate(yearmonth = substring(day, 0, 7)) %>%
  pull(yearmonth)





```

## Data visualisation





- Boxplots of the hourly concentrations non-voc





- Boxplots sulfur compounds, NOx, ozone



#### • Boxplots VOCs







## Data pre-processing

- STEP 1: Limits of detection

```
# Define LOD for each chemical
LOD_non_voc <- c('ch4' = 0.9,
                 'co2' = 0.0433,
                 'co' = 40,
                 'h2s' = 0.4,
                 'so2' = 0.4,
                 'nox' = 0.05,
                 'o3' = 1)

# LOD_voc_monthly <- read_csv('../data/LNM_VOC_LOD_Rounded.csv') %>% select(-1)
#
```

```

# # extract the yearmonth from date variables
# LOD_voc_monthly <- LOD_voc_monthly %>%
#   mutate(yearmonth = strftime(as.POSIXct(start_date, format = '%Y-%m-%d %H:%M:%S',
#                                           tz = 'UTC'), '%Y-%m'))
#
# LOD_voc_monthly <- LOD_voc_monthly %>%
#   select(-c(start_date, end_date)) %>%
#   select(!any_of(ends_with('half_ldl')))
#
# colnames(LOD_voc_monthly) <- str_replace_all(names(LOD_voc_monthly), '_ldl', '')

LOD_voc_avg <- read_xlsx('../data/LNM_VOC_Uncertainties.xlsx', skip = 1)
LOD_voc_avg <- LOD_voc_avg %>%
  select(1, 4) %>%
  rename('LOD' = 2, 'chemical' = 1) %>%
  head(20)

```

- STEP 2: Background correction

##	ch4	co2	co	h2s	so2
##	1928.000	411.300	59.910	0.200	0.200
##	nox	o3	ethane	ethene	propane
##	0.025	0.500	0.916	0.011	0.224
##	propene	1_3-butadiene	i-butane	n-butane	acetylene
##	0.009	0.007	0.035	0.090	0.019
##	cyclopentane	i-pentane	n-pentane	n-hexane	isoprene
##	0.005	0.038	0.042	0.021	0.005
##	n-heptane	benzene	n-octane	toluene	ethyl-benzene
##	0.004	0.017	0.004	0.004	0.004
##	m&p-xylene	o-xylene			
##	0.004	0.004			

- Summary statistics of backgrounds and extremes

```

get_info <- function(column) {
  N <- length(column)
  background <- quantile(column, 0)
  quantile1 <- quantile(column, 0.01)
  quantile99 <- quantile(column, 0.99)
  n_background <- sum(column == background)
  max <- max(column)
  return(c(N, quantile1, quantile99, max, background, n_background))
}

info_table <- hourly_full_nona %>%
  reframe(across(everything(), ~ get_info(.x)))

info_table <- info_table %>%
  mutate(rownames = c('N', '1st percentile', '99th percentile', 'Max',
                      'Background', '# Background')) %>%
  pivot_longer(-rownames) %>%
  pivot_wider(names_from = rownames, values_from = value)

knitr::kable(info_table)

```

name	N	1st percentile	99th percentile	Max	Background	Background	#
ch4	4788	1962.98700	6286.12400	34010.900	1928.000		1
co2	4788	416.47870	460.62260	503.990	411.300		1
co	4788	84.23050	442.08860	2513.440	59.910		1
h2s	4788	0.20000	5.20986	27.700	0.200		829
so2	4788	0.20000	1.78686	8.578	0.200		3266
nox	4788	0.22974	89.72371	452.959	0.025		2
o3	4788	0.50000	76.02600	103.100	0.500		259
ethane	4788	1.84422	526.44700	2060.000	0.916		1
ethene	4788	0.01100	3.50826	16.970	0.011		163
propane	4788	0.84674	300.79000	1211.000	0.224		1
propene	4788	0.00900	0.69739	5.528	0.009		411
1_3-butadiene	4788	0.00700	0.05900	1.207	0.007		3357
i-butane	4788	0.15148	60.89400	296.600	0.035		1
n-butane	4788	0.37248	166.52100	536.900	0.090		1
acetylene	4788	0.04900	2.61304	8.471	0.019		2
cyclopentane	4788	0.00500	3.06899	13.460	0.005		96
i-pentane	4788	0.10987	49.60210	215.900	0.038		1
n-pentane	4788	0.10487	55.95980	258.800	0.042		1
n-hexane	4788	0.04300	18.17780	93.360	0.021		2
isoprene	4788	0.00500	0.03313	0.362	0.005		2816
n-heptane	4788	0.01500	6.57669	30.470	0.004		5
benzene	4788	0.02800	3.78693	9.610	0.017		3
n-octane	4788	0.00400	2.00839	6.867	0.004		100
toluene	4788	0.01300	3.52165	9.077	0.004		11
ethyl-benzene	4788	0.00400	0.31613	0.931	0.004		918
m&p-xylene	4788	0.00400	1.29156	3.123	0.004		851
o-xylene	4788	0.00400	0.45700	0.922	0.004		1330

- STEP 2 processing continued: background correction
- adjustments that were made according to paper: Gunnar's paper section 2.2 and Guha 3.3
- Check whether chemical has background noise level that needs to be removed
- NO ADJUSTMENT if minimum value < 2xLOD and maximum value > 100xLOD

```
adjusting_neg_bg_from_lod <- function(chemical, LOD, background, hourly_data){
  # get min and max
  min_value <- min(hourly_data[chemical], na.rm = TRUE)
  max_value <- max(hourly_data[chemical], na.rm = TRUE)
  # if min less than double LOD or max > 100 times LOD
  # adjust to -100 (for entire column???)
  if (min_value < 2 * LOD & max_value > 100 * LOD ){
    return (0)
  }
  return (background)
}
```

- Check if background is negligible for non voc
- merge background and LOD

```
background_lod_non_voc <- tibble(chemical = non_vocs,
                                LOD = LOD_non_voc,
                                background = unname(background_levels[non_vocs]))
adjusted_background_non_voc <- background_lod_non_voc %>%
```

```

rowwise() %>%
mutate(min = min(hourly_full_nona[chemical], na.rm = TRUE),
       LODx2 = 2 * LOD,
       criterion1 = min(hourly_full_nona[chemical], na.rm = TRUE) < 2 * LOD,
       max = max(hourly_full_nona[chemical], na.rm = TRUE),
       LODx100 = 100 * LOD,
       criterion2 = max(hourly_full_nona[chemical], na.rm = TRUE) > 100 * LOD,
       adjusted_background = adjusting_neg_bg_from_lod(chemical, LOD, background,
                                                    hourly_full_nona))

```

- Check if background is negligible for voc
- merge background and LOD

```

background_lod_voc <- LOD_voc_avg %>%
  left_join(tibble(chemical = setdiff(names(background_levels), non_vocs),
                  background = background_levels[setdiff(names(background_levels),
                                                         non_vocs)]))

adjusted_background_voc <- background_lod_voc %>%
  rowwise() %>%
  mutate(min = min(hourly_full_nona[chemical], na.rm = TRUE),
         LODx2 = 2 * LOD,
         criterion1 = min(hourly_full_nona[chemical], na.rm = TRUE) < 2 * LOD,
         max = max(hourly_full_nona[chemical], na.rm = TRUE),
         LODx100 = 100 * LOD,
         criterion2 = max(hourly_full_nona[chemical], na.rm = TRUE) > 100 * LOD,
         adjusted_background = adjusting_neg_bg_from_lod(chemical, LOD, background,
                                                         hourly_full_nona))

```

- create dataset with background removed

```

# So now we have the adjusted background concentrations
hourly_nona_bgrm <- hourly_full_nona %>%
  mutate(across(adjusted_background_non_voc$chemical,
               ~ .x - adjusted_background_non_voc$adjusted_background[
                 adjusted_background_non_voc$chemical == cur_column()])))
hourly_nona_bgrm <- hourly_nona_bgrm %>%
  mutate(across(adjusted_background_voc$chemical,
               ~ .x - adjusted_background_voc$adjusted_background[
                 adjusted_background_voc$chemical == cur_column()])))

```

- check number of 0 values per compound

```

# look at zero values
colSums(hourly_nona_bgrm == 0)

```

```

##          ch4          co2          co          h2s          so2
##          1          1          1          829          3266
##          nox          o3          ethane          ethene          propane
##          0          0          1          0          1
##          propene 1_3-butadiene          i-butane          n-butane          acetylene
##          0          3357          1          1          0
##          cyclopentane          i-pentane          n-pentane          n-hexane          isoprene
##          0          1          1          2          2816
##          n-heptane          benzene          n-octane          toluene          ethyl-benzene
##          0          0          0          0          0
##          m&p-xylene          o-xylene

```



```
##           0           0
```

- STEP 3: replace zero values with random values between 0 and 0.5xLOD

```
set.seed(123)
replace_zero_with_random <- function(column, name, LOD_df){
  LOD <- LOD_df$LOD[LOD_df$chemical == name]
  column <- if_else(column == 0, round(runif(length(column), 0, 0.5 * LOD), 3), column)
  return (column)
}

hourly_nona_bgrm_zerorepl <- hourly_nona_bgrm %>%
  mutate(across(adjusted_background_non_voc$chemical,
    ~ replace_zero_with_random(.x, cur_column(), adjusted_background_non_voc)))

hourly_nona_bgrm_zerorepl <- hourly_nona_bgrm_zerorepl %>%
  mutate(across(adjusted_background_voc$chemical,
    ~ replace_zero_with_random(.x, cur_column(), adjusted_background_voc)))
```

- STEP 4: Normalize the non-vocs

```
#normalizing function
normalize_column <- function(column){
  background <- quantile(column, 0)
  max <- quantile(column, 1) # this could be adjusted
  return ((column - background)/(max - background))
}
```

- STEP 4: Normalize all

```
# normalize all
hourly_nona_bgrm_zerorepl_norm <- as_tibble(sapply(as.list(hourly_nona_bgrm_zerorepl),
  normalize_column))

#normalize the NON_VOC
summary(hourly_nona_bgrm_zerorepl_norm)
```

```
##           ch4           co2           co           h2s
## Min.      :0.00000   Min.      :0.0000   Min.      :0.00000   Min.      :0.00000
## 1st Qu.:0.00579   1st Qu.:0.1384   1st Qu.:0.02592   1st Qu.:0.01022
## Median :0.01460   Median :0.1823   Median :0.03884   Median :0.02335
## Mean     :0.02683   Mean     :0.2000   Mean     :0.04761   Mean     :0.03501
## 3rd Qu.:0.03720   3rd Qu.:0.2418   3rd Qu.:0.05970   3rd Qu.:0.04525
## Max.     :1.00000   Max.     :1.0000   Max.     :1.00000   Max.     :1.00000
##           so2           nox           o3           ethane
## Min.      :0.000000   Min.      :0.000000   Min.      :0.00000   Min.      :0.000000
## 1st Qu.:0.007878   1st Qu.:0.006534   1st Qu.:0.09747   1st Qu.:0.008385
## Median :0.015994   Median :0.020262   Median :0.25487   Median :0.026671
## Mean     :0.026287   Mean     :0.036440   Mean     :0.26676   Mean     :0.050992
## 3rd Qu.:0.023633   3rd Qu.:0.049978   3rd Qu.:0.40546   3rd Qu.:0.075375
## Max.     :1.000000   Max.     :1.000000   Max.     :1.00000   Max.     :1.000000
##           ethene           propane           propene           1_3-butadiene
## Min.      :0.00000   Min.      :0.000000   Min.      :0.000000   Min.      :0.000000
## 1st Qu.:0.01268   1st Qu.:0.009283   1st Qu.:0.005979   1st Qu.:0.002500
## Median :0.03547   Median :0.028409   Median :0.018482   Median :0.004167
## Mean     :0.05042   Mean     :0.053803   Mean     :0.028772   Mean     :0.007371
## 3rd Qu.:0.07266   3rd Qu.:0.080130   3rd Qu.:0.042761   3rd Qu.:0.007500
## Max.     :1.00000   Max.     :1.000000   Max.     :1.000000   Max.     :1.000000
```

##	i-butane	n-butane	acetylene	cyclopentane
##	Min. :0.00000	Min. :0.000000	Min. :0.00000	Min. :0.000000
##	1st Qu.:0.00614	1st Qu.:0.008777	1st Qu.:0.02674	1st Qu.:0.007432
##	Median :0.01925	Median :0.027522	Median :0.05135	Median :0.022668
##	Mean :0.03837	Mean :0.054900	Mean :0.07436	Mean :0.043730
##	3rd Qu.:0.05369	3rd Qu.:0.077042	3rd Qu.:0.10211	3rd Qu.:0.062653
##	Max. :1.00000	Max. :1.000000	Max. :1.00000	Max. :1.000000
##	i-pentane	n-pentane	n-hexane	isoprene
##	Min. :0.000000	Min. :0.000000	Min. :0.000000	Min. :0.000000
##	1st Qu.:0.006303	1st Qu.:0.005681	1st Qu.:0.004703	1st Qu.:0.002801
##	Median :0.019941	Median :0.018371	Median :0.016039	Median :0.005602
##	Mean :0.041094	Mean :0.038859	Mean :0.034979	Mean :0.010315
##	3rd Qu.:0.057857	3rd Qu.:0.054837	3rd Qu.:0.049544	3rd Qu.:0.011204
##	Max. :1.000000	Max. :1.000000	Max. :1.000000	Max. :1.000000
##	n-heptane	benzene	n-octane	toluene
##	Min. :0.000000	Min. :0.00000	Min. :0.000000	Min. :0.00000
##	1st Qu.:0.005473	1st Qu.:0.01637	1st Qu.:0.008269	1st Qu.:0.01389
##	Median :0.018348	Median :0.04222	Median :0.026009	Median :0.04276
##	Mean :0.039328	Mean :0.07655	Mean :0.054341	Mean :0.07825
##	3rd Qu.:0.055866	3rd Qu.:0.10779	3rd Qu.:0.076497	3rd Qu.:0.11333
##	Max. :1.000000	Max. :1.00000	Max. :1.000000	Max. :1.00000
##	ethyl-benzene	m&p-xylene	o-xylene	
##	Min. :0.000000	Min. :0.000000	Min. :0.00000	
##	1st Qu.:0.007551	1st Qu.:0.007374	1st Qu.:0.00000	
##	Median :0.034520	Median :0.039115	Median :0.04139	
##	Mean :0.062378	Mean :0.077508	Mean :0.08650	
##	3rd Qu.:0.090615	3rd Qu.:0.115742	3rd Qu.:0.12881	
##	Max. :1.000000	Max. :1.000000	Max. :1.00000	

- FINAL step: create matrix of processed and normalized concentrations for NMF

```
normalized_matrix <- as.matrix(hourly_nona_bgrm_zerorepl_norm)
#important: using the normalized VOCs for this file
```

## NMF section

### Preprocess

#### Global variables

```
components <- 4:10
```

#### Remove Ozone

```
normalized_matrix_less_o3 <- normalized_matrix[, setdiff(colnames(normalized_matrix), "o3")]
```

#### Compute error matrix

```
# compute weight matrix (uncertainties or error)
# Based on the Guha paper

weight_matrix <- matrix(0, nrow = nrow(normalized_matrix_less_o3),
                        ncol = ncol(normalized_matrix_less_o3))
LOD_merged <- tibble(chemical = c(adjusted_background_non_voc$chemical,
```

```

                                adjusted_background_voc$chemical),
                                LOD = c(adjusted_background_non_voc$LOD,
                                adjusted_background_voc$LOD))

LOD_merged <- tibble(chemical = names(hourly_nona_bgrm_zerorepl_norm)) %>%
  left_join(LOD_merged) %>%
  filter(chemical %in% colnames(normalized_matrix_less_o3))

## Joining with `by = join_by(chemical)`
# creating uncertainty Matrix
for (i in 1:dim(weight_matrix)[1]) {
  for (j in 1:dim(weight_matrix)[2]) {
    chemical <- colnames(normalized_matrix_less_o3)[j]
    xij <- normalized_matrix_less_o3[i, j]
    LOD <- LOD_merged$LOD[LOD_merged$chemical == chemical]
    # Get LOD value for this row
    if (j == 1) {
      # based on equation 6, we sqrt ch4 (at column = 1) and times by 1
      weight_matrix[i, j] <- sqrt(xij)
    } else if (j == 2) {
      # 0.25 for co2
      weight_matrix[i, j] <- 0.25 * sqrt(xij)
    } else if (j == 3) {
      # 0.5 for CO
      weight_matrix[i, j] <- 0.5 * sqrt(xij)
    } else if (xij <= LOD) {
      weight_matrix[i, j] <- 2 * LOD # equation 5a) in reference paper
    } else {
      weight_matrix[i, j] <- sqrt(((0.1 * xij)**2 + LOD**2)) #equation 5c) in reference paper
    }
  }
}

```

Helper for source contributions plots

## LS-NMF + Random seed


```

# for each rank, run 30 models and find best one
# start_time_lsnmf_rand <- Sys.time()
#
# lsnmf_random_less_o3 <- nmf(
#   normalized_matrix_less_o3,
#   components,
#   method = "ls-nmf",
#   weight = weight_matrix,
#   30,
#   seed = 123456
# )
#
# end_time_lsnmf_rand <- Sys.time()
# end_time_lsnmf_rand-start_time_lsnmf_rand
# # 19.25 minutes to run the above
#

```

```
# saveRDS(lsnmf_random_less_o3,
#         'lsnmf_random_less_o3.rds')
lsnmf_random_less_o3 <- readRDS('lsnmf_random_less_o3.rds')
```

```
# plots the NMF rank survey
plot(lsnmf_random_less_o3)
```



Look at 5 factors:

```
output <- lsnmf_random_less_o3$fit$`5`
W <- basis(output)
H <- coef(output)
```

```
# Convert H to a data frame for ggplot
H_df_5c_less_o3 <- as.data.frame(H)
# Add a column for chemicals
H_df_5c_less_o3$Component <- rownames(H_df_5c_less_o3)

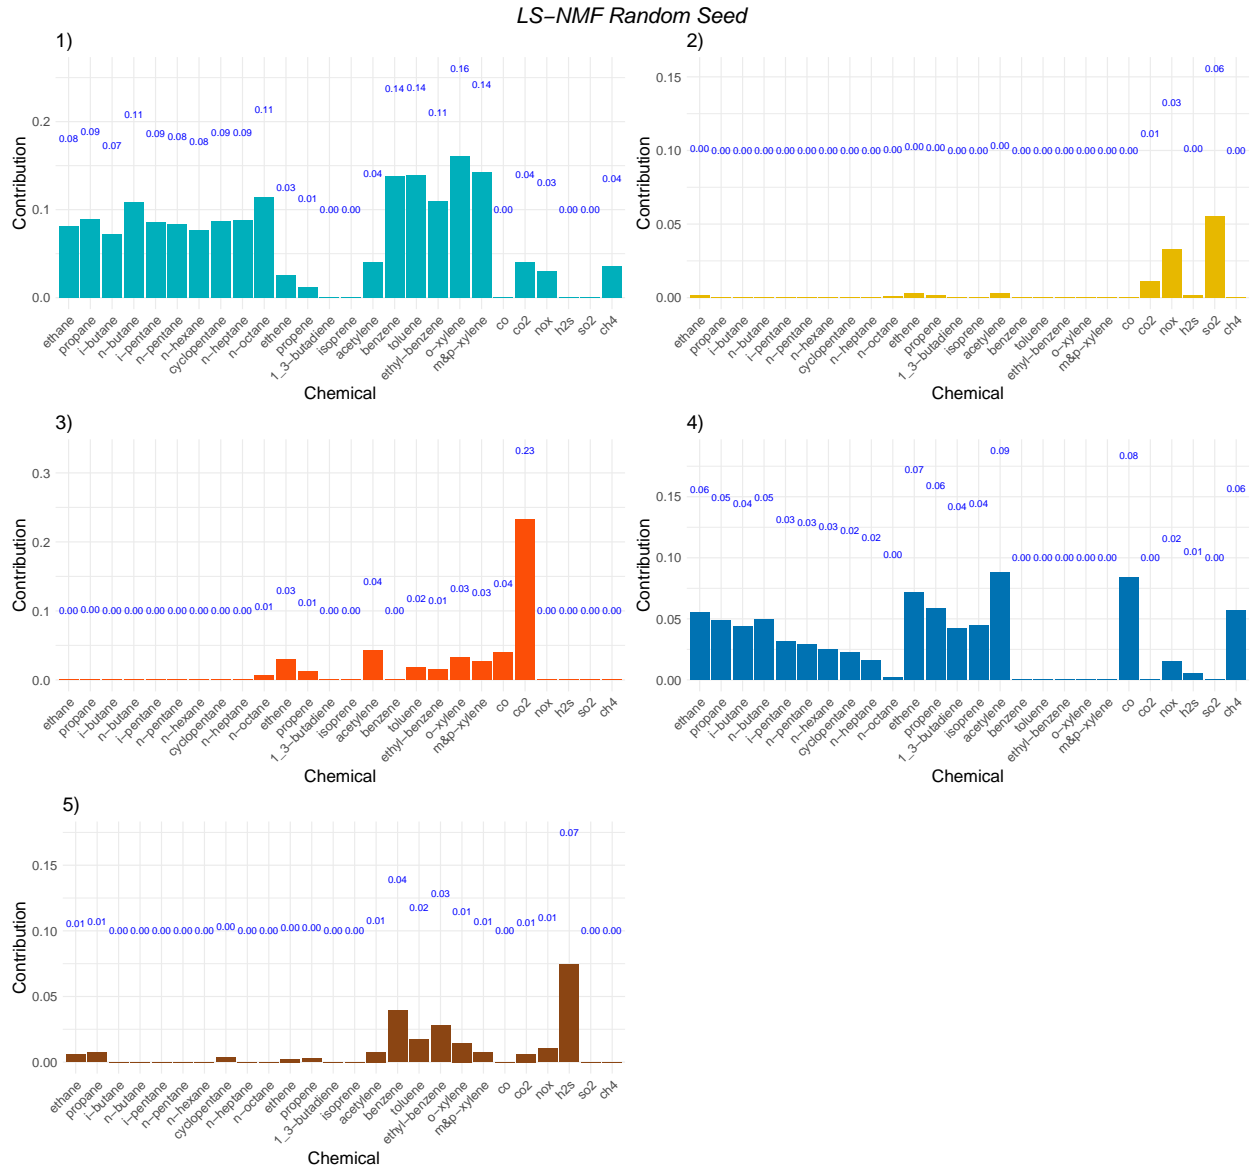
# Reshape data to long format
H_long_5c_less_o3 <- pivot_longer(H_df_5c_less_o3, cols = -Component,
                                   names_to = "Chemical", values_to = "Contribution")

# Plot
nmfplt_1_lsnmf_random_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                                         '1', '1')
nmfplt_2_lsnmf_random_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
```

```

                                '2', '2')
nmfplt_3_lsnmf_random_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                '3', '3')
nmfplt_4_lsnmf_random_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                '4', '4')
nmfplt_5_lsnmf_random_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                '5', '5')

```



## LS-NMF + nndsvd seed

```

#function below used to estimate the optimal rank and will be used in the nmf() function.
# start_time_lsnmf_nndsvd <- Sys.time()
#
# lsnmf_nndsvd_less_o3 <- nmf(

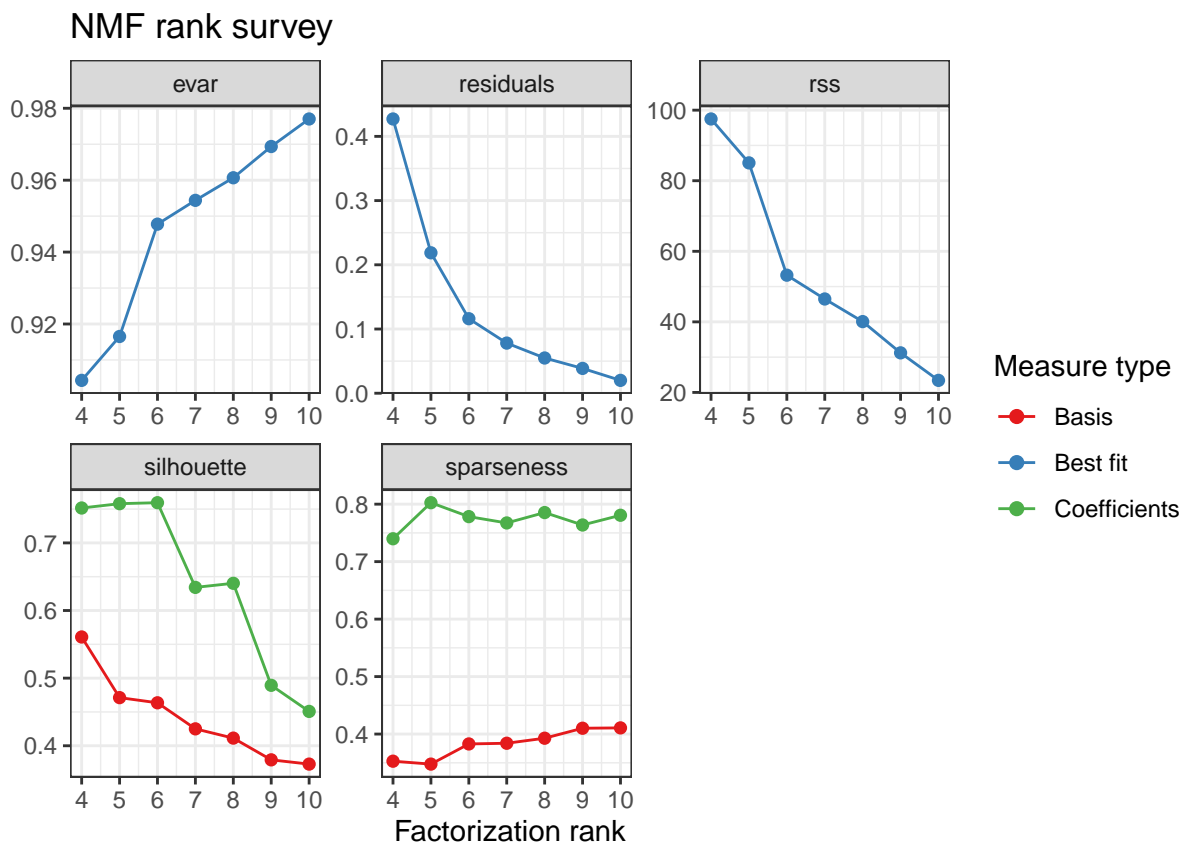
```

```

# normalized_matrix_less_o3,
# rank = components,
# nrun = 1, # since using nndsvd
# method = "ls-nmf",
# weight = weight_matrix,
# seed = 'nndsvd'
# )
#
# end_time_lsnmf_nndsvd <- Sys.time()
# end_time_lsnmf_nndsvd - start_time_lsnmf_nndsvd
# # 1.34 minutes to run the above
#
# saveRDS(lsnmf_nndsvd_less_o3,
#         'lsnmf_nndsvd_less_o3.rds')
lsnmf_nndsvd_less_o3 <- readRDS('lsnmf_nndsvd_less_o3.rds')

# plots the NMF rank survey
plot(lsnmf_nndsvd_less_o3)

```



Look at 5 factors:

```

output <- lsnmf_nndsvd_less_o3$fit$`5`
W <- basis(output)
H <- coef(output)

```

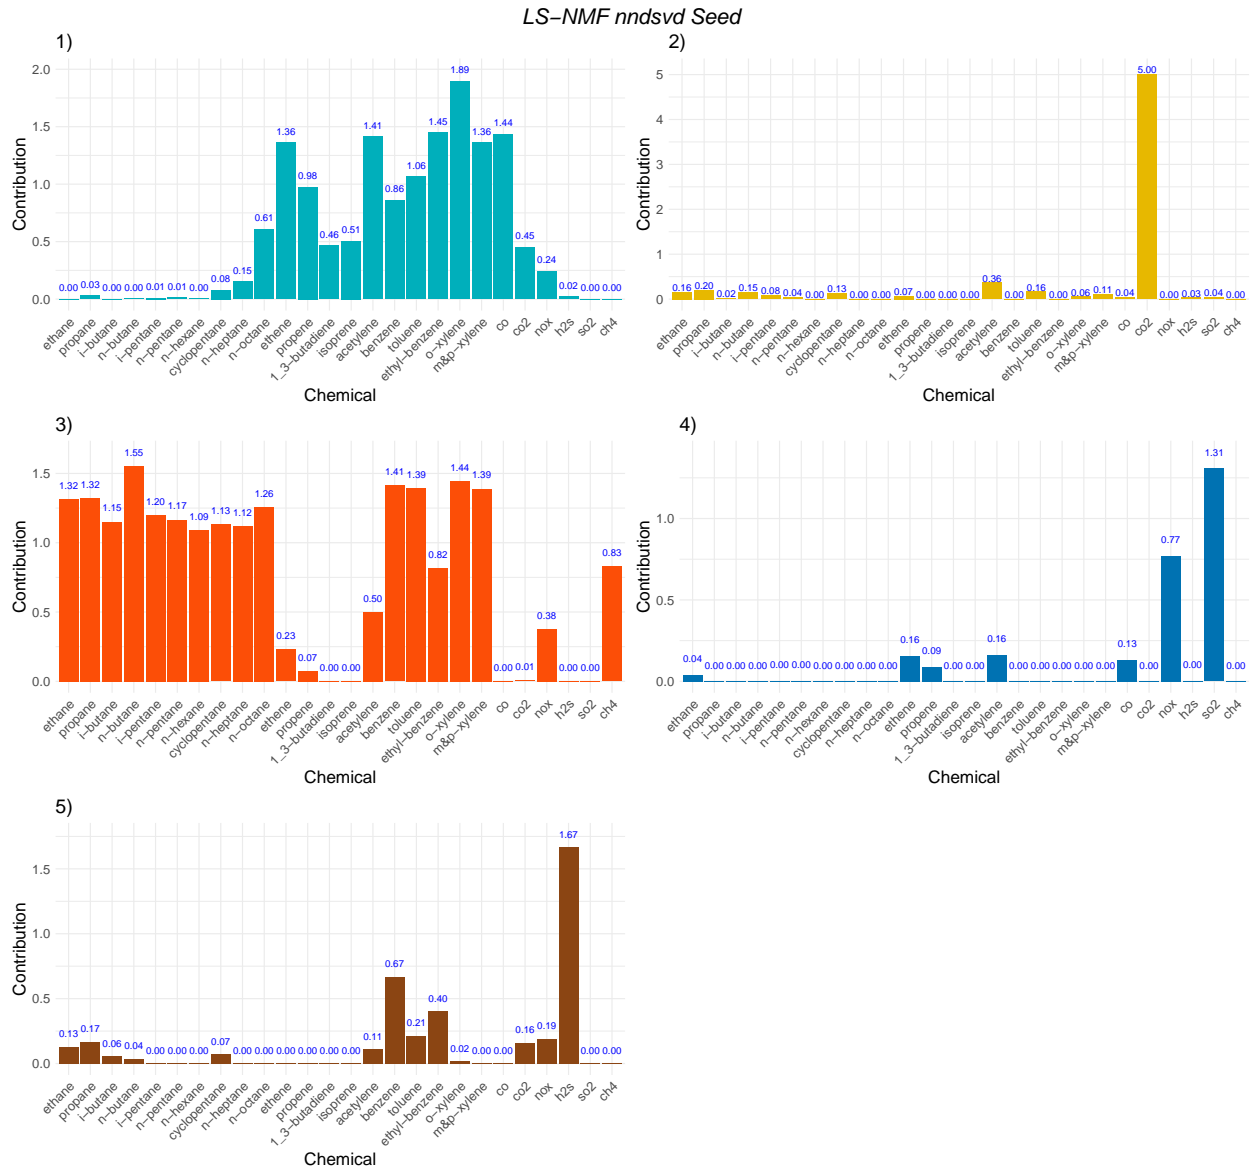
```

# Convert H to a data frame for ggplot
H_df_5c_less_o3 <- as.data.frame(H)
# Add a column for chemicals
H_df_5c_less_o3$Component <- rownames(H_df_5c_less_o3)

# Reshape data to long format
H_long_5c_less_o3 <- pivot_longer(H_df_5c_less_o3, cols = -Component,
                                   names_to = "Chemical", values_to = "Contribution")

# Plot
nmfplt_1_lsnmf_nndsvd_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                                         '1', '1')
nmfplt_2_lsnmf_nndsvd_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                                         '2', '2')
nmfplt_3_lsnmf_nndsvd_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                                         '3', '3')
nmfplt_4_lsnmf_nndsvd_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                                         '4', '4')
nmfplt_5_lsnmf_nndsvd_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                                         '5', '5')

```




## Comparing random seed vs nndsvd for ls-nmf

Residuals is defined as  $\sum((X - \text{fitted}(\text{object})) * \text{weight})^2 / 2$

```
tibble(component = components,
  random_residual = lsnmf_random_less_o3$measures$residuals,
  nndsvd_residual = lsnmf_nndsvd_less_o3$measures$residuals) %>%
  ggplot() +
  geom_line(aes(x = component, y = random_residual, color = 'Random seed')) +
  geom_point(aes(x = component, y = random_residual, color = 'Random seed')) +
  geom_line(aes(x = component, y = nndsvd_residual, color = 'nndsvd')) +
  geom_point(aes(x = component, y = nndsvd_residual, color = 'nndsvd')) +
  scale_colour_manual("",
    breaks = c("Random seed", "nndsvd"),
    values = c("tomato1", "dodgerblue")) +
  labs(x = 'Rank', y = 'Residual') +
  theme_bw()
```





## KL + Random seed


```
# start_time_kl_random <- Sys.time()
#
# kl_random_less_o3 <- nmf(
#   normalized_matrix_less_o3,
#   rank = components,
#   nrun = 30,
#   method = "KL",
#   seed = 123456
# )
#
# end_time_kl_random <- Sys.time()
# end_time_kl_random - start_time_kl_random
# 14.27 minutes to run the above

# saveRDS(kl_random_less_o3, 'kl_random_less_o3.rds')

kl_random_less_o3 <- readRDS('kl_random_less_o3.rds')

# plots the NMF rank survey
plot(kl_random_less_o3)
```

## NMF rank survey



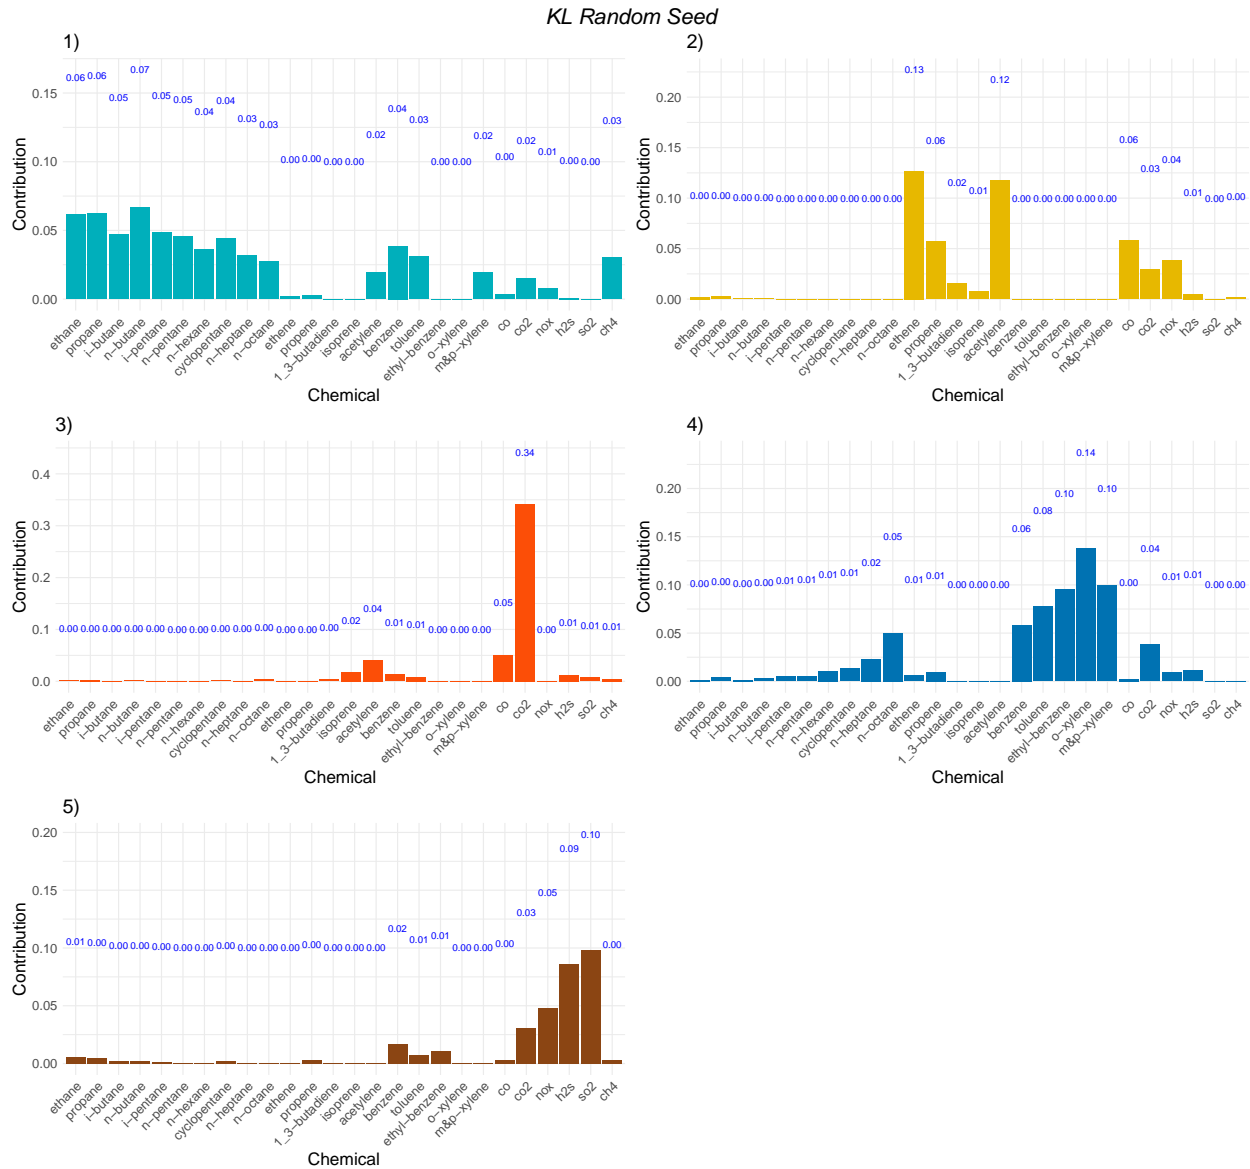
### Look at 5 factors:

```
output <- kl_random_less_o3$fit$`5`
W <- basis(output)
H <- coef(output)
```

```
# Convert H to a data frame for ggplot
H_df_5c_less_o3 <- as.data.frame(H)
# Add a column for chemicals
H_df_5c_less_o3$Component <- rownames(H_df_5c_less_o3)

# Reshape data to long format
H_long_5c_less_o3 <- pivot_longer(H_df_5c_less_o3, cols = -Component,
                                   names_to = "Chemical", values_to = "Contribution")

# Plot
nmfplt_1_kl_random_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                                    '1', '1')
nmfplt_2_kl_random_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                                    '2', '2')
nmfplt_3_kl_random_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                                    '3', '3')
nmfplt_4_kl_random_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                                    '4', '4')
nmfplt_5_kl_random_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                                    '5', '5')
```



## KL + nndsvd

```
# errors <- numeric(length(components) - 4)

# Loop over the number of components
# for (n in components) {
#   nmf_result <- nmf(normalized_matrix_less_o3, rank = n, method = "KL", seed='nndsvd')
#   reconstruction <- basis(nmf_result) %*% coef(nmf_result)
#   error <- norm(normalized_matrix_less_o3 - reconstruction, type = "F")^2 # RSS
#   errors[n-3] <- error
#   print(paste0('Completed ', n - 3, ' out of 7'))
# }
#
# saveRDS(errors, 'errors_KL_nndsvd_less_o3.rds')
#
# errors <- readRDS('errors_KL_nndsvd_less_o3.rds')
```

```

# start_time_kl_nndsvd <- Sys.time()
#
# kl_nndsvd_less_o3 <- nmf(
#   normalized_matrix_less_o3,
#   rank = components,
#   nrun = 1,
#   method = "KL",
#   seed = 'nndsvd'
# )
#
# end_time_kl_nndsvd <- Sys.time()
# end_time_kl_nndsvd - start_time_kl_nndsvd
# 1 minute to run the above

# saveRDS(kl_nndsvd_less_o3, 'kl_nndsvd_less_o3.rds')


kl_nndsvd_less_o3 <- readRDS('kl_nndsvd_less_o3.rds')

```

```

# plots the NMF rank survey
plot(kl_nndsvd_less_o3)

```



Look at 5 factors:

```

output <- kl_nndsvd_less_o3$fit$`5`
W <- basis(output)

```

```

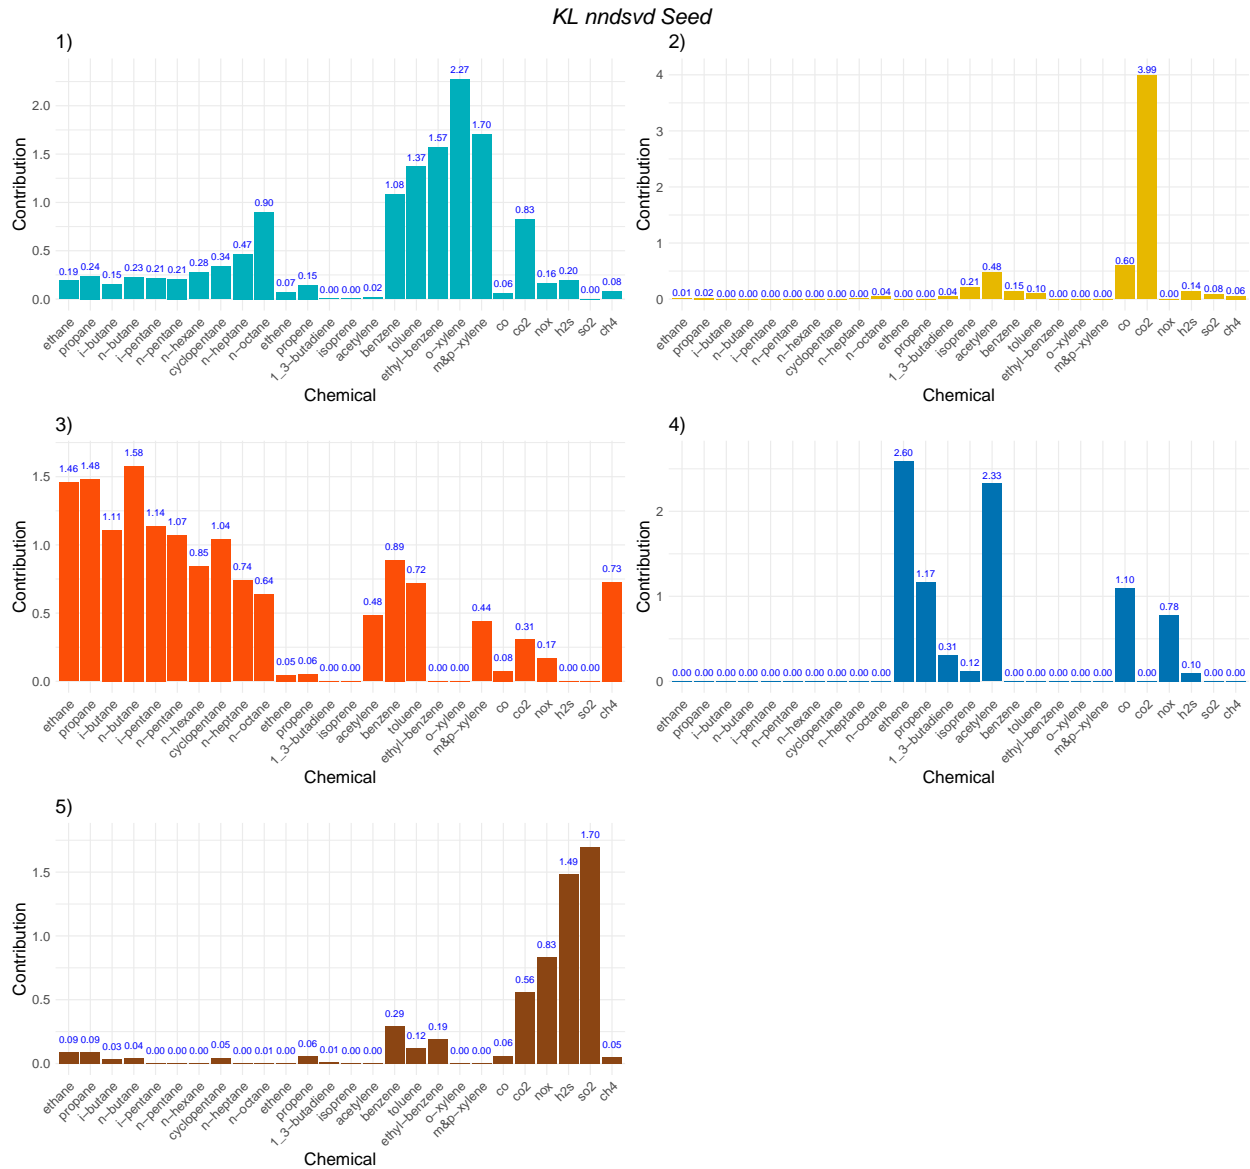
H <- coef(output)

# Convert H to a data frame for ggplot
H_df_5c_less_o3 <- as.data.frame(H)
# Add a column for chemicals
H_df_5c_less_o3$Component <- rownames(H_df_5c_less_o3)

# Reshape data to long format
H_long_5c_less_o3 <- pivot_longer(H_df_5c_less_o3, cols = -Component,
                                   names_to = "Chemical", values_to = "Contribution")

# Plot
nmfplt_1_kl_nndsvd_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                                    '1', '1')
nmfplt_2_kl_nndsvd_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                                    '2', '2')
nmfplt_3_kl_nndsvd_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                                    '3', '3')
nmfplt_4_kl_nndsvd_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                                    '4', '4')
nmfplt_5_kl_nndsvd_less_o3_5c <- get_component_plot(H_long_5c_less_o3,
                                                    '5', '5')

```



## Comparing all four methods

Using RSS, WRSS, and KL

```
get_residual <- function(component, seed, method, objective) {
  fitted <- fitted(get(paste(method, seed, 'less_o3', sep = '_'))$fit[[component-3]])
  if (objective == 'wrss') {
    return(sum(((normalized_matrix_less_o3 - fitted) * weight_matrix)^2)/2)
  } else if (objective == 'kl') {
    log_term <- normalized_matrix_less_o3/fitted
    log_term[log_term<.Machine$double.eps] <- .Machine$double.eps
    return(sum(normalized_matrix_less_o3 * log(log_term) + fitted - normalized_matrix_less_o3))
  } else if (objective == 'rss') {
    return(norm(normalized_matrix_less_o3 - fitted, type = 'F')^2)
  }
}
```

```

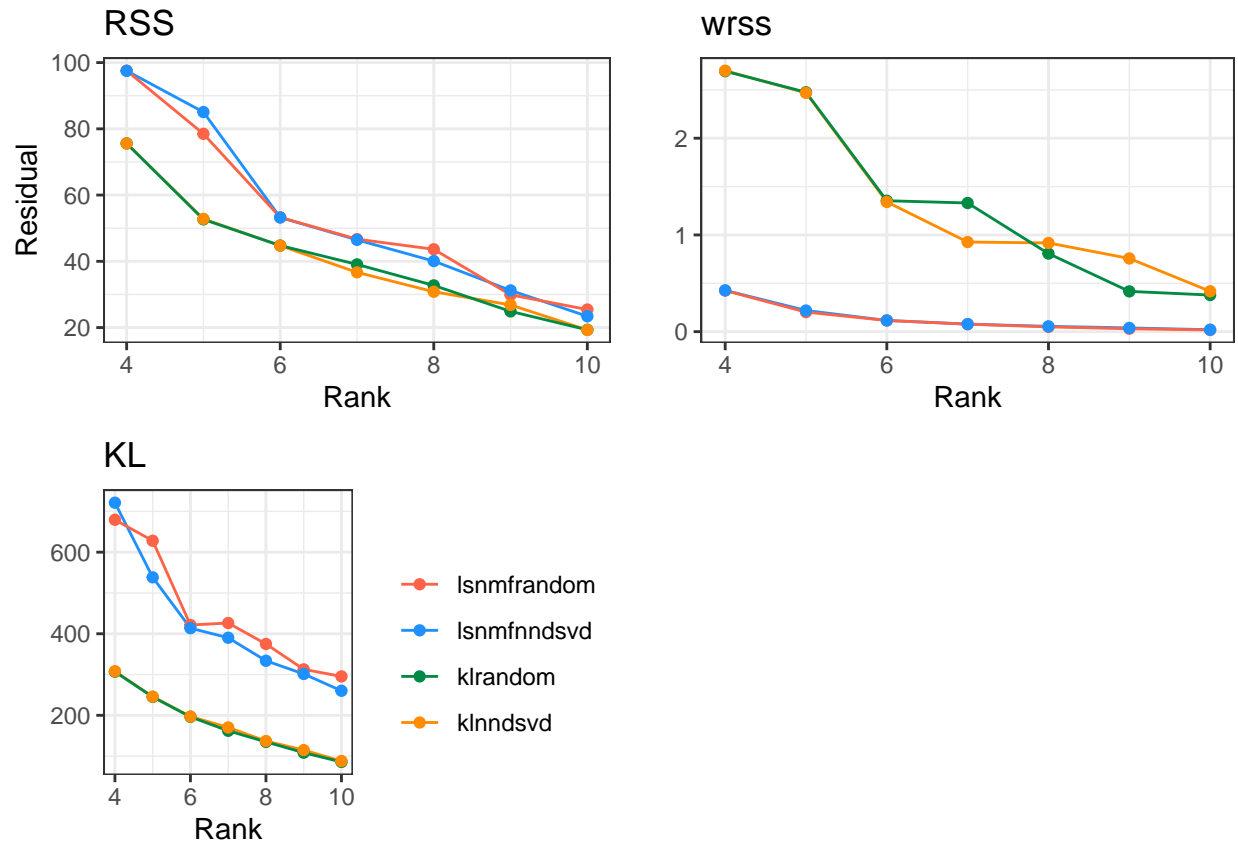
df_plot <- expand_grid(
  component = components,
  seed = c('random', 'nndsvd'),
  method = c('lsnmf', 'kl'),
  objective = c('rss', 'wrss', 'kl')
) %>%
  rowwise() %>%
  mutate(residual = get_residual(component, seed, method, objective)) %>%
  ungroup() %>%
  mutate(model = paste0(method, seed))

RSS_plot <- df_plot %>%
  filter(objective=='rss') %>%
  ggplot() +
  geom_line(aes(x = component, y = residual, group = model, color = model)) +
  geom_point(aes(x = component, y = residual, group = model, color = model)) +
  scale_colour_manual("",
    breaks = c("lsnmfrandom", "lsnmfnndsvd",
               "klrandom", "klndsvd"),
    values = c("tomato1", "dodgerblue",
               "springgreen4", "darkorange")) +
  labs(x = 'Rank', y = 'Residual', title = 'RSS') +
  theme_bw() +
  theme(legend.position="none")

WRSS_plot <- df_plot %>%
  filter(objective=='wrss') %>%
  ggplot() +
  geom_line(aes(x = component, y = residual, group = model, color = model)) +
  geom_point(aes(x = component, y = residual, group = model, color = model)) +
  scale_colour_manual("",
    breaks = c("lsnmfrandom", "lsnmfnndsvd",
               "klrandom", "klndsvd"),
    values = c("tomato1", "dodgerblue",
               "springgreen4", "darkorange")) +
  labs(x = 'Rank', y = '', title = 'wrss') +
  theme_bw() +
  theme(legend.position="none")

KL_plot <- df_plot %>%
  filter(objective=='kl') %>%
  ggplot() +
  geom_line(aes(x = component, y = residual, group = model, color = model)) +
  geom_point(aes(x = component, y = residual, group = model, color = model)) +
  scale_colour_manual("",
    breaks = c("lsnmfrandom", "lsnmfnndsvd",
               "klrandom", "klndsvd"),
    values = c("tomato1", "dodgerblue",
               "springgreen4", "darkorange")) +
  labs(x = 'Rank', y = '', title = 'KL') +
  theme_bw()
grid.arrange(RSS_plot, WRSS_plot, KL_plot, ncol=2)

```



## NMF with 5 source factors without ozone

- remove ozone
- use KL divergence loss with svd seed
- Extract W (basis) and H (coefs) matrices
- Calculate variance explained in all 5 factors
- Calculate variance explained by each factor

```
nmf_result_5c_less_o3 <- kl_nndsvd_less_o3$fit$`5`

basis_matrix_5c_less_o3 <- basis(nmf_result_5c_less_o3) #W
coef_matrix_5c_less_o3 <- coef(nmf_result_5c_less_o3) #H

# get variance explained by the factors (total residuals)
reconstruct<-fitted(nmf_result_5c_less_o3)

tss <- sum((normalized_matrix_less_o3 - mean(normalized_matrix_less_o3))^2)
rss <- sum((normalized_matrix_less_o3 - reconstruct)^2)
variance_explained <- 1 - (rss / tss)
variance_explained
```

```
## [1] 0.9212864
```

```
# get variance explained by each factor separately
# Compute variance explained by each factor
# Initialize variance explained tracker
variance_explained_factors <- numeric(5)
```



```

# Incrementally add factors and calculate variance explained
reconstruction <- matrix(0, nrow = nrow(basis_matrix_5c_less_o3), ncol = ncol(coef_matrix_5c_less_o3))

for (i in 1:5) {
  # Add the i-th factor to the reconstruction
  reconstruction <- reconstruction + (basis_matrix_5c_less_o3[, i, drop=FALSE] %*% coef_matrix_5c_less_o3[, i, drop=FALSE])

  # Compute Residual Sum of Squares (RSS)
  rss_f <- sum((normalized_matrix_less_o3 - reconstruction)^2)

  # Compute Variance Explained by adding this factor
  variance_explained_factors[i] <- 1 - (rss_f / tss)
}

# Print variance explained by each factor cumulatively
variance_explained_factors


```

```
## [1] 0.2395401 0.5113683 0.8113445 0.8921360 0.9212864
```

```

par(mfrow = c(1, 2))
image(basis_matrix_5c_less_o3, main = "Basis Matrix (W)")
image(coef_matrix_5c_less_o3, main = "Coefficient Matrix (H)")

```



```

# Convert H to a data frame for ggplot
H_df_5c_less_o3 <- as.data.frame(coef_matrix_5c_less_o3)
# Add a column for chemicals

```


```

H_df_5c_less_o3$Component <- rownames(H_df_5c_less_o3)

# Reshape data to long format
H_long_5c_less_o3 <- pivot_longer(H_df_5c_less_o3, cols = -Component,
                                   names_to = "Chemical", values_to = "Contribution")

# Plot
nmfplt_1_svd_5c_less_o3 <- get_component_plot(H_long_5c_less_o3,
                                              '1', '1) Traffic emissions factor')
nmfplt_2_svd_5c_less_o3 <- get_component_plot(H_long_5c_less_o3,
                                              '2', '2) Combustion from engines, turbines, compressors')
nmfplt_3_svd_5c_less_o3 <- get_component_plot(H_long_5c_less_o3,
                                              '3', '3) Oil & gas fugitive and venting emissions')
nmfplt_4_svd_5c_less_o3 <- get_component_plot(H_long_5c_less_o3,
                                              '4', '4) Flaring, incomplete combustion')
nmfplt_5_svd_5c_less_o3 <- get_component_plot(H_long_5c_less_o3,
                                              '5', '5) Flaring, sour gas')

```



## Fingerprint plot

```
custom_colors <- c("Traffic" = "#00AFBB",
  "Combustion" = "#E7B800",
  "Oil & Gas Emissions" = "#FC4E07",
  "Flare (factor 4)" = "#0072B2",
  "Flare (factor 5)" = "#8B4513")


contrib_prop <- apply(H_df_5c_less_o3[,1:(length(H_df_5c_less_o3)-1)], MARGIN = 2,
  FUN = function(x) {x/sum(x)})

contrib_prop %>%
  as_tibble() %>%
  mutate(
    Component = c(
      'Traffic',
```

```

      'Combustion',
      'Oil & Gas Emissions',
      'Flare (factor 4)',
      'Flare (factor 5)'
    )
  ) %>%
mutate(Component = factor(
  Component,
  levels = c(
    'Traffic',
    'Combustion',
    'Oil & Gas Emissions',
    'Flare (factor 4)',
    'Flare (factor 5)'
  )
)) %>%
pivot_longer(cols = -Component,
  names_to = "Chemical",
  values_to = "Contribution_prop") %>%
mutate(Chemical = factor(Chemical, levels = desired_order)) %>%
ggplot(aes(fill = Component, y = Contribution_prop, x = Chemical)) +
geom_bar(position = "fill", stat = "identity") +
scale_fill_manual(values = custom_colors) +
theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
labs(x = "Chemical", y = "Contribution Proportion") +
theme(
  panel.grid.major = element_blank(),
  panel.grid.minor = element_blank(),
  panel.background = element_blank()
)

```



```
#ggsave("fingerprint.png", c)
```

## Wind plots

```
hourly_wind_nona <- hourly_nona %>%
  select(wdr_deg, wsp_ms)

data_to_plot <- tibble(
  component1 = basis(nmf_result_5c_less_o3)[,1],
  component2 = basis(nmf_result_5c_less_o3)[,2],
  component3 = basis(nmf_result_5c_less_o3)[,3],
  component4 = basis(nmf_result_5c_less_o3)[,4],
  component5 = basis(nmf_result_5c_less_o3)[,5],
  wd = round(hourly_wind_nona$wdr_deg, -1)
)

color_pal <- c("#00AFBB", "#E7B800", "#FC4E07", "#0072B2", "#8B4513")

data_long <- data_to_plot %>%
  pivot_longer(cols = starts_with("component"), names_to = "Factor", values_to = "Expression")

factor_labels <- c(
  "component1" = "Factor 1 - Traffic",
  "component2" = "Factor 2 - Combustion",
```

```

"component3" = "Factor 3 - Oil & Gas Emissions",
"component4" = "Factor 4 - Flaring",
"component5" = "Factor 5 - Flaring Sour Gas"
)

data_long <- data_long %>%
  mutate(wd = factor(wd, levels = sort(unique(wd))))

# Select every second wind direction for labeling
every_second_label <- levels(data_long$wd)[seq(1, length(levels(data_long$wd)), by = 2)]

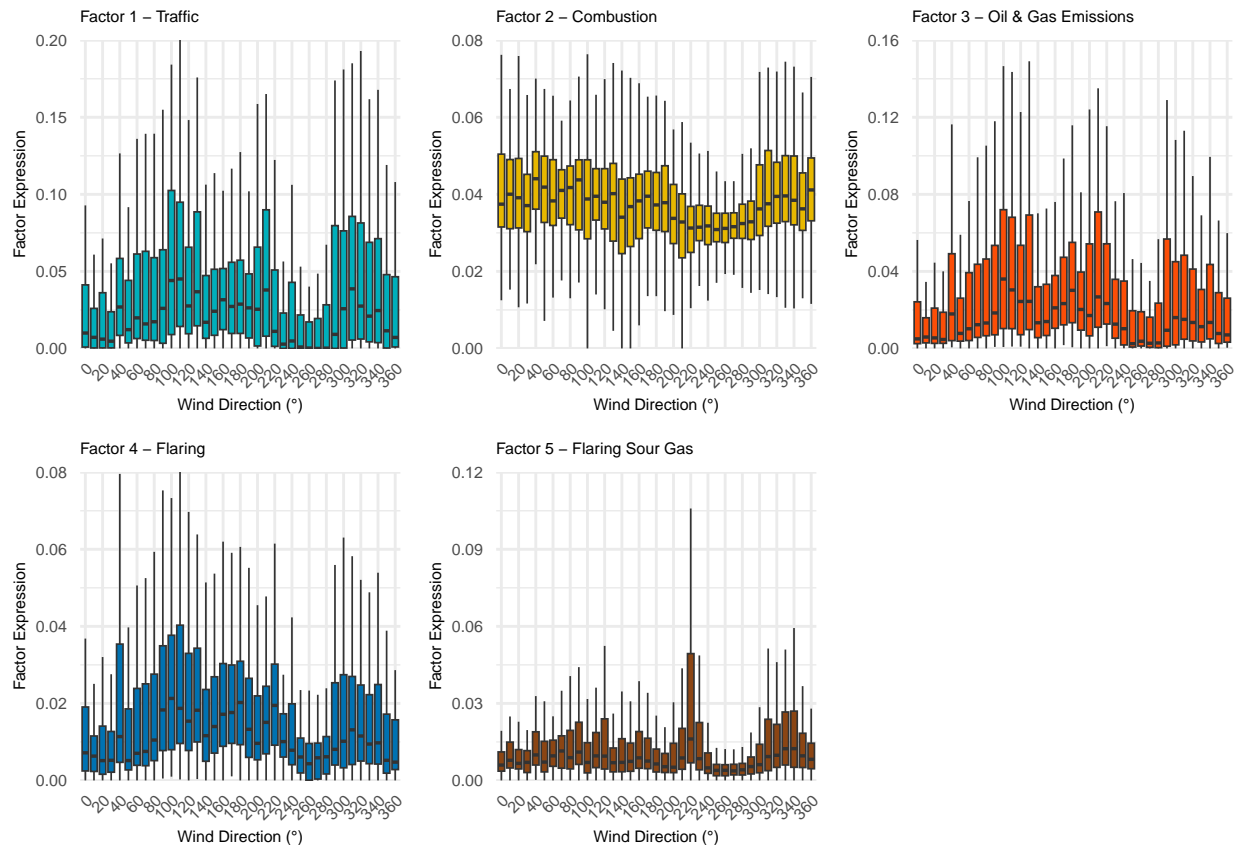
y_axis_limits <- list(
  "component1" = c(0, 0.2),
  "component2" = c(0, 0.08),
  "component3" = c(0, 0.16),
  "component4" = c(0, 0.08),
  "component5" = c(0, 0.12)
)

plots <- lapply(1:5, function(i) {
  factor_name <- paste0("component", i)

  ggplot(data_long %>% filter(Factor == factor_name),
    aes(x = wd, y = Expression, fill = as.factor(wd))) +
    geom_boxplot(outliers=F, size=0.3) +
    scale_fill_manual(values = rep(color_pal[i], length(unique(data_long$wd)))) +
    scale_x_discrete(breaks = every_second_label) +
    coord_cartesian(ylim = y_axis_limits[[factor_name]]) +
    scale_y_continuous(
      limits = c(0, NA),
      breaks = seq(0, y_axis_limits[[factor_name]][2], length.out = 5) ,
      expand=expansion(mult=c(0))
    ) +
    labs(title = factor_labels[factor_name],
      x = "Wind Direction (°)",
      y = "Factor Expression") +
    theme_minimal() +
    theme(
      legend.position = "none",
      plot.title = element_text(size = 6), # Smaller title text
      axis.title = element_text(size = 6), # Smaller axis labels
      axis.text = element_text(size = 6), # Smaller x and y tick labels
      axis.text.x = element_text(angle = 45, hjust = 1)
    )
})

grid.arrange(grobs = plots, ncol = 3)

```



```
# ggsave("factors-wind.png", w)
```

## Factor analysis

- merge in factors 1-5 to dataset (hourly)

```
# First look at how well this approximates
fitted_5c_less_o3 <- fitted(nmf_result_5c_less_o3)
sum(abs(normalized_matrix_less_o3-fitted_5c_less_o3))

## [1] 1059.63

# NMF factorizes  $V = WH$ 
# Store Basis matrix (W) and Coef Matrix (H)
saveRDS(basis_matrix_5c_less_o3, 'result_rfiles/nmf_norm_5c_less_o3_basis.rds')
saveRDS(coef_matrix_5c_less_o3, 'result_rfiles/nmf_norm_5c_less_o3_coef.rds')

# Merge basis matrix into hourly observations
basis_matrix_5c_less_o3 <- as_tibble(basis_matrix_5c_less_o3) %>%
  setNames(c('Factor1', 'Factor2', 'Factor3', 'Factor4', 'Factor5'))
```

```
## Warning: The `x` argument of `as_tibble.matrix()` must have unique column names if
## `.name_repair` is omitted as of tibble 2.0.0.
## i Using compatibility `.name_repair`.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.
```

```
normalized_hourly_data_5c_less_o3 <- hourly_nona[,c('day', 'time_utc')] %>%
  cbind(normalized_matrix_less_o3) %>%
  cbind(basis_matrix_5c_less_o3) %>%
  right_join(hourly_data %>% select(-'day'), join_by(time_utc), suffix = c('_norm', ''))

# saveRDS(normalized_hourly_data_5c_less_o3,
# 'result_rfiles/normalized_hourly_data_5c_less_o3.rds')
normalized_hourly_data_5c_less_o3 <- readRDS('result_rfiles/normalized_hourly_data_5c_less_o3.rds')
```

- make daily dataset for VNF analysis
- compute wind directions from plots

```
# Also compute a daily dataset
normalized_daily_data_5c_less_o3 <- normalized_hourly_data_5c_less_o3 %>%
  group_by(day) %>%
  summarise(across(where(is.numeric) & !any_of('wdr_deg'), ~ mean(.x, na.rm = T)),
    wdr_deg = as.numeric(mean(circular(wdr_deg, units = "degrees"), na.rm = T))) %>%
  mutate(wdr_deg = if_else(wdr_deg < 0, wdr_deg+360, wdr_deg)) %>%
  mutate(wind_45_135 = wdr_deg >= 45 & wdr_deg < 135,
    wind_135_180 = wdr_deg >= 135 & wdr_deg < 180,
    wind_180_270 = wdr_deg >= 180 & wdr_deg < 270,
    wind_270_45 = wdr_deg >= 270 & wdr_deg < 45)

# saveRDS(normalized_daily_data_5c_less_o3,
# 'result_rfiles/normalized_daily_data_5c_less_o3.rds')

normalized_daily_data_5c_less_o3 <-
  readRDS('result_rfiles/normalized_daily_data_5c_less_o3.rds')
```

- 1) number of flares in 100km of trailer associated with NMF
- 2) weighted count based on distance to trailer

```
# Check if relationship between # flares and flare factor (4 & 5)
# Linear model
flare_factor <- lm(n_flare_100 ~ Factor1 + Factor2 + Factor3 + Factor4 + Factor5,
  data = normalized_daily_data_5c_less_o3)
summary(flare_factor)
```

```
##
## Call:
## lm(formula = n_flare_100 ~ Factor1 + Factor2 + Factor3 + Factor4 +
##   Factor5, data = normalized_daily_data_5c_less_o3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -54.638 -22.160   4.205  18.488  76.270
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    52.438     7.455   7.034 1.61e-11 ***
## Factor1       -27.402    106.172  -0.258  0.7965
## Factor2      -338.560    196.573  -1.722  0.0861 .
## Factor3       286.534    151.310   1.894  0.0593 .
## Factor4      -287.536    244.717  -1.175  0.2410
## Factor5       231.978    212.510   1.092  0.2760
```



```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 27.79 on 273 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared:  0.03877,    Adjusted R-squared:  0.02117
## F-statistic: 2.202 on 5 and 273 DF,  p-value: 0.05434

flare_factor45 <- lm(n_flare_100 ~ Factor4 + Factor5, data = normalized_daily_data_5c_less_o3)
summary(flare_factor45)

##
## Call:
## lm(formula = n_flare_100 ~ Factor4 + Factor5, data = normalized_daily_data_5c_less_o3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -53.409 -23.830   5.588  18.235  77.131
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   40.780      3.386  12.044 <2e-16 ***
## Factor4       -48.431     150.383  -0.322  0.7477
## Factor5        360.559     206.393   1.747  0.0818 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 28.02 on 276 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared:  0.01171,    Adjusted R-squared:  0.004548
## F-statistic: 1.635 on 2 and 276 DF,  p-value: 0.1968

flare_factor_weighted <- lm(weighted.count ~ Factor1 + Factor2 + Factor3 + Factor4 + Factor5,
                             data = normalized_daily_data_5c_less_o3)
summary(flare_factor_weighted)

##
## Call:
## lm(formula = weighted.count ~ Factor1 + Factor2 + Factor3 + Factor4 +
##      Factor5, data = normalized_daily_data_5c_less_o3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.3295 -0.2180  0.0546  0.3809  3.9848
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    2.2572     0.2278   9.907 <2e-16 ***
## Factor1         0.1244     3.2450   0.038  0.9694
## Factor2        -4.7740     6.0080  -0.795  0.4275
## Factor3         7.4339     4.6246   1.607  0.1091
## Factor4       -12.9155     7.4794  -1.727  0.0853 .
## Factor5         4.0762     6.4951   0.628  0.5308
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 0.8492 on 273 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared: 0.02221, Adjusted R-squared: 0.0043
## F-statistic: 1.24 on 5 and 273 DF, p-value: 0.2905

flare_factor_weighted45 <- lm(weighted.count ~ Factor4 + Factor5,
                             data = normalized_daily_data_5c_less_o3)
summary(flare_factor_weighted45)

##
## Call:
## lm(formula = weighted.count ~ Factor4 + Factor5, data = normalized_daily_data_5c_less_o3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2558 -0.1821  0.0775  0.3622  3.9366
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   2.0910     0.1029  20.315 <2e-16 ***
## Factor4       -4.0144     4.5712  -0.878  0.381
## Factor5        7.3663     6.2738   1.174  0.241
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8518 on 276 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared: 0.005595, Adjusted R-squared: -0.001611
## F-statistic: 0.7765 on 2 and 276 DF, p-value: 0.461

# All factors + wind speed + wind direction + factor5:sw wind.
# Wind direction from 270 to 45 is left as reference group.
flare_factor_weighted_2 <- lm(weighted.count ~ Factor1 + Factor2 + Factor3 +
                             Factor4 + Factor5 + wsp_ms + wind_45_135 +
                             wind_135_180 + Factor5*wind_180_270,
                             data = normalized_daily_data_5c_less_o3)
summary(flare_factor_weighted_2)

##
## Call:
## lm(formula = weighted.count ~ Factor1 + Factor2 + Factor3 + Factor4 +
##      Factor5 + wsp_ms + wind_45_135 + wind_135_180 + Factor5 *
##      wind_180_270, data = normalized_daily_data_5c_less_o3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.4382 -0.2036  0.0782  0.3578  3.9528
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   2.14418     0.33390   6.422 6.11e-10 ***
## Factor1        0.07365     3.35780   0.022  0.9825
## Factor2       -4.92773     6.19462  -0.795  0.4270
## Factor3        8.99603     4.80767   1.871  0.0624 .
```

```

## Factor4                -10.66714    7.82960   -1.362    0.1742
## Factor5                 3.28556    7.52447    0.437    0.6627
## wsp_ms                  0.04430    0.04523    0.980    0.3282
## wind_45_135TRUE        -0.15557    0.17743   -0.877    0.3814
## wind_135_180TRUE       -0.15153    0.13270   -1.142    0.2545
## wind_180_270TRUE       -0.21667    0.22628   -0.958    0.3392
## Factor5:wind_180_270TRUE 2.51821   13.08348    0.192    0.8475
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8524 on 268 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared:  0.03302,    Adjusted R-squared:  -0.003061
## F-statistic: 0.9152 on 10 and 268 DF,  p-value: 0.5196

# Same as above but only factor 4 and 5
flare_factor_weighted_3 <- lm(weighted.count ~ Factor4 + Factor5 + wsp_ms +
                             Factor5*wind_180_270,
                             data = normalized_daily_data_5c_less_o3)
summary(flare_factor_weighted_3)

##
## Call:
## lm(formula = weighted.count ~ Factor4 + Factor5 + wsp_ms + Factor5 *
##     wind_180_270, data = normalized_daily_data_5c_less_o3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2706 -0.1961  0.0714  0.3721  3.9358
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      2.057728   0.222033   9.268  <2e-16 ***
## Factor4          -3.273445   5.155494  -0.635   0.526
## Factor5           7.499668   7.297235   1.028   0.305
## wsp_ms           0.009441   0.040855   0.231   0.817
## wind_180_270TRUE -0.059345   0.215638  -0.275   0.783
## Factor5:wind_180_270TRUE -0.078349  12.897092  -0.006   0.995
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.856 on 273 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared:  0.006554,    Adjusted R-squared:  -0.01164
## F-statistic: 0.3602 on 5 and 273 DF,  p-value: 0.8754

# Same as above but interaction between factor 4 and SW wind
flare_factor_weighted_3b <- lm(weighted.count ~ Factor4 + Factor5 + wsp_ms +
                              Factor4*wind_180_270,
                              data = normalized_daily_data_5c_less_o3)
summary(flare_factor_weighted_3b)

##
## Call:
## lm(formula = weighted.count ~ Factor4 + Factor5 + wsp_ms + Factor4 *
##     wind_180_270, data = normalized_daily_data_5c_less_o3)

```

```
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.3058 -0.2123  0.0650  0.3774  3.9523
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      2.030311    0.215853   9.406  <2e-16 ***
## Factor4          -0.841684    5.604781  -0.150    0.881
## Factor5           7.401203    6.365967   1.163    0.246
## wsp_ms           0.005478    0.040870   0.134    0.893
## wind_180_270TRUE  0.143565    0.224432   0.640    0.523
## Factor4:wind_180_270TRUE -10.510472    9.618074  -1.093    0.275
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8541 on 273 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared:  0.01088,    Adjusted R-squared:  -0.007235
## F-statistic: 0.6006 on 5 and 273 DF,  p-value: 0.6995

# Same as above but with East wind
flare_factor_weighted_3c <- lm(weighted.count ~ Factor4 + Factor5 + wsp_ms +
                             Factor5*wind_45_135,
                             data = normalized_daily_data_5c_less_o3)
summary(flare_factor_weighted_3c)

##
## Call:
## lm(formula = weighted.count ~ Factor4 + Factor5 + wsp_ms + Factor5 *
##     wind_45_135, data = normalized_daily_data_5c_less_o3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2355 -0.1832  0.0761  0.3768  3.9129
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      2.1177298    0.2207356   9.594  <2e-16 ***
## Factor4          -4.0986004    5.1120629  -0.802    0.423
## Factor5           6.1846170    6.5540446   0.944    0.346
## wsp_ms           0.0008777    0.0407876   0.022    0.983
## wind_45_135TRUE  -0.4275903    0.3752245  -1.140    0.255
## Factor5:wind_45_135TRUE 22.0187173    23.3341439   0.944    0.346
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8543 on 273 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared:  0.01053,    Adjusted R-squared:  -0.007588
## F-statistic: 0.5813 on 5 and 273 DF,  p-value: 0.7143

flare_factor_weighted_3d <- lm(weighted.count ~ Factor4 + Factor5 + wsp_ms +
                             Factor4*wind_45_135,
                             data = normalized_daily_data_5c_less_o3)
```

```
summary(flare_factor_weighted_3d)
```

```
##
## Call:
## lm(formula = weighted.count ~ Factor4 + Factor5 + wsp_ms + Factor4 *
##     wind_45_135, data = normalized_daily_data_5c_less_o3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2680 -0.1822  0.0707  0.3665  3.9260
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      2.09473    0.21997   9.523  <2e-16 ***
## Factor4          -4.50558    5.36882  -0.839    0.402
## Factor5           7.78444    6.38818   1.219    0.224
## wsp_ms            0.00343    0.04074   0.084    0.933
## wind_45_135TRUE  -0.19324    0.28449  -0.679    0.498
## Factor4:wind_45_135TRUE  4.84677   13.38860   0.362    0.718
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8555 on 273 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared:  0.007783, Adjusted R-squared: -0.01039
## F-statistic: 0.4283 on 5 and 273 DF, p-value: 0.8288
# Wind speed + factor 4 and interaction with East wind
flare_factor_weighted_4a <- lm(weighted.count ~ wsp_ms + Factor4*wind_45_135,
                              data = normalized_daily_data_5c_less_o3)
summary(flare_factor_weighted_4a)
```

```
##
## Call:
## lm(formula = weighted.count ~ wsp_ms + Factor4 * wind_45_135,
##     data = normalized_daily_data_5c_less_o3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.1668 -0.1983  0.0661  0.3882  3.8663
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      2.187226    0.206642  10.585  <2e-16 ***
## wsp_ms          -0.003785    0.040343  -0.094    0.925
## Factor4         -2.428778    5.095632  -0.477    0.634
## wind_45_135TRUE  -0.171010    0.284161  -0.602    0.548
## Factor4:wind_45_135TRUE  3.879541   13.376874   0.290    0.772
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8562 on 274 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared:  0.002386, Adjusted R-squared: -0.01218
```

```
## F-statistic: 0.1638 on 4 and 274 DF, p-value: 0.9565
# Wind speed + factor 4 and interaction with SE wind
flare_factor_weighted_4b <- lm(weighted.count ~ wsp_ms + Factor4*wind_135_180,
                               data = normalized_daily_data_5c_less_o3)
summary(flare_factor_weighted_4b)

##
## Call:
## lm(formula = weighted.count ~ wsp_ms + Factor4 * wind_135_180,
##     data = normalized_daily_data_5c_less_o3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2269 -0.2186  0.0794  0.3693  3.8404
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      2.229232   0.202818  10.991  <2e-16 ***
## wsp_ms           0.001297   0.039954   0.032   0.9741
## Factor4          -5.716640   5.512131  -1.037   0.3006
## wind_135_180TRUE -0.440258   0.229503  -1.918   0.0561 .
## Factor4:wind_135_180TRUE 18.354002   9.757795   1.881   0.0610 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.851 on 274 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared:  0.01456, Adjusted R-squared:  0.0001782
## F-statistic: 1.012 on 4 and 274 DF, p-value: 0.4014
# Wind speed + factor 4 and interaction with SW wind
flare_factor_weighted_4c <- lm(weighted.count ~ wsp_ms + Factor4*wind_180_270,
                               data = normalized_daily_data_5c_less_o3)
summary(flare_factor_weighted_4c)

##
## Call:
## lm(formula = weighted.count ~ wsp_ms + Factor4 * wind_180_270,
##     data = normalized_daily_data_5c_less_o3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.1911 -0.1952  0.0539  0.4087  3.8937
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      2.121372   0.201270  10.540  <2e-16 ***
## wsp_ms          -0.001567   0.040444  -0.039   0.969
## Factor4           1.058349   5.364665   0.197   0.844
## wind_180_270TRUE  0.143445   0.224576   0.639   0.524
## Factor4:wind_180_270TRUE -10.634150   9.623656  -1.105   0.270
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8547 on 274 degrees of freedom
```

```
## (1 observation deleted due to missingness)
## Multiple R-squared: 0.005983, Adjusted R-squared: -0.008528
## F-statistic: 0.4123 on 4 and 274 DF, p-value: 0.7997

# Wind speed + factor 5 and interaction with East wind
flare_factor_weighted_5a <- lm(weighted.count ~ wsp_ms + Factor5*wind_45_135,
                               data = normalized_daily_data_5c_less_o3)
summary(flare_factor_weighted_5a)

##
## Call:
## lm(formula = weighted.count ~ wsp_ms + Factor5 * wind_45_135,
##     data = normalized_daily_data_5c_less_o3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2299 -0.1901  0.0789  0.3725  3.9421
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      2.01637    0.18083   11.151  <2e-16 ***
## wsp_ms           0.01530    0.03658    0.418    0.676
## Factor5          4.62551    6.25483    0.740    0.460
## wind_45_135TRUE -0.40644    0.37405   -1.087    0.278
## Factor5:wind_45_135TRUE 21.30900   23.30214    0.914    0.361
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8537 on 274 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared: 0.008204, Adjusted R-squared: -0.006275
## F-statistic: 0.5666 on 4 and 274 DF, p-value: 0.6871

# Wind speed + factor 5 and interaction with SE wind
flare_factor_weighted_5b <- lm(weighted.count ~ wsp_ms + Factor5*wind_135_180,
                               data = normalized_daily_data_5c_less_o3)
summary(flare_factor_weighted_5b)

##
## Call:
## lm(formula = weighted.count ~ wsp_ms + Factor5 * wind_135_180,
##     data = normalized_daily_data_5c_less_o3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2259 -0.1876  0.0693  0.3775  3.9290
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      2.02295    0.18889   10.710  <2e-16 ***
## wsp_ms           0.01822    0.03643    0.500    0.618
## Factor5          4.19811    7.25463    0.579    0.563
## wind_135_180TRUE -0.15866    0.20317   -0.781    0.436
## Factor5:wind_135_180TRUE 6.06208   12.22399    0.496    0.620
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 0.8545 on 274 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared: 0.006345, Adjusted R-squared: -0.008161
## F-statistic: 0.4374 on 4 and 274 DF, p-value: 0.7816

# Wind speed + factor 5 and interaction with SW wind
flare_factor_weighted_5c <- lm(weighted.count ~ wsp_ms + Factor5*wind_180_270,
                              data = normalized_daily_data_5c_less_o3)
summary(flare_factor_weighted_5c)

##
## Call:
## lm(formula = weighted.count ~ wsp_ms + Factor5 * wind_180_270,
##     data = normalized_daily_data_5c_less_o3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2696 -0.2009  0.0671  0.3748  3.9556
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      1.98021     0.18525  10.689  <2e-16 ***
## wsp_ms           0.02098     0.03655   0.574   0.566
## Factor5          6.25646     7.02199   0.891   0.374
## wind_180_270TRUE -0.06960     0.21480  -0.324   0.746
## Factor5:wind_180_270TRUE -0.16687    12.88229  -0.013   0.990
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8551 on 274 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared: 0.005087, Adjusted R-squared: -0.009438
## F-statistic: 0.3502 on 4 and 274 DF, p-value: 0.8438

# Check relationship between avg flare distance and flare factor (4 & 5)
# Linear model
flare_factor_dist <- lm(distToLovi ~ Factor4 + Factor5, data = normalized_daily_data_5c_less_o3)
summary(flare_factor_dist)

##
## Call:
## lm(formula = distToLovi ~ Factor4 + Factor5, data = normalized_daily_data_5c_less_o3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -34.449  -2.443  -0.139   2.266  31.399
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   54.7821     0.8756  62.564  <2e-16 ***
## Factor4        7.2656     39.7289   0.183   0.855
## Factor5       64.5797     52.2590   1.236   0.218
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```



```
##  
## Residual standard error: 6.883 on 252 degrees of freedom  
## (25 observations deleted due to missingness)  
## Multiple R-squared: 0.008425, Adjusted R-squared: 0.0005557  
## F-statistic: 1.071 on 2 and 252 DF, p-value: 0.3444
```